

Importing necessary libraries

```
In [1]: from annoy import AnnoyIndex
import numpy as np
import pandas as pd
import re
import cohere
```

```
In [2]: API_KEY = "52ibaZZU3WoKzduJrpEI2eNhyEikplQ3QTLdD9BP"
co = cohere.Client(API_KEY)
```

important data

```
In [3]: text = """This is a fascinating time in the study and application of large language models. New advancements are
being made in the field of natural language processing, and the integration of large language models into various applications
is becoming increasingly common. In this guide, I share my analysis of the current architectural best practices for data-informed language model
development, focusing on the integration of external data sources and the use of retrieval-augmented generation (RAG) techniques.

Overview
In nearly all practical applications of large language models (LLM's), there are instances in which you want the model to
reference specific data. This data can be used to provide context, answer specific questions, or generate content based on
external information. At a high level, there are two primary methods for referencing specific data:

1. Context-based approach: Insert data as context in the model prompt, and direct the response to utilize that information.
2. Fine-tuning approach: Fine-tune a model, by providing hundreds or thousands of prompt <=> completion pairs.

Shortcomings of Knowledge Retrieval for Existing LLM's
Both of these methods have significant shortcomings in isolation.

For the context-based approach:
Models have a limited context size, with the latest `davinci-003` model only able to process up to 4,000 tokens.
Processing more tokens equates to longer processing times. In customer-facing scenarios, this impairs the user experience.
Processing more tokens equates to higher API costs, and may not lead to more accurate responses if the information is
irrelevant or noisy.

For the fine-tuning approach:
Generating prompt <=> completion pairs is time-consuming and potentially expensive.
Many repositories from which you want to reference information are quite large. For example, if your application requires
access to a vast corpus of text, fine-tuning a model to reference that data can be prohibitively expensive.
Some external data sources change quickly. For example, it is not optimal to retrain a customer support model based on
outdated information.
Best practices around fine-tuning are still being developed. LLM's themselves can be used to assist with the generation of
prompt-completion pairs.

The Solution, Simplified
The design above goes by various names, most commonly "retrieval-augmented generation" or "RETRO". Links & references:
- RAG: Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks
- RETRO: Improving language models by retrieving from trillions of tokens
- REALM: Retrieval-Augmented Language Model Pre-Training
- Retrieval-augmented generation a) retrieves relevant data from outside of the language model (non-parametric) and b)
generates a response based on that data.

Retrieval
The retrieval of relevant information is worth further explanation. As you can see, data may come from multiple sources.
Language model embeddings are numerical representations of concepts in text and seem to have endless uses. Here, we use
them to find relevant information in a corpus of text.

Back to the flow – when a user submits a question, an LLM processes the message in multiple ways, but the key step is
retrieving relevant information. This is where the RAG technique comes in.

Augmentation
Building the prompt using the relevant text chunks is straightforward. The prompt begins with some basic prompt
template, and the relevant text chunks are inserted into the prompt. Finally, the combined prompt is sent to the large language model. An answer is parsed from the completion and provided to the user.

That's it! While this is a simple version of the design, it's inexpensive, accurate, and perfect for many lightweight
applications.

Advanced Design
I want to take a moment to discuss several research developments that may enter into the retrieval-augmented generation
design.

Generate-then-Read Pipelines
This category of approaches involves processing the user input with an LLM before retrieving relevant data.
Basically, a user's question lacks some of the relevance patterns that an informative answer will display. For example,
a user might ask "What is the capital of France?" and the LLM might respond "The capital of France is Paris." A similar
approach titled generate-then-read (GenRead) builds on the practice by implementing a clustering algorithm to identify
relevant information in a corpus of text.

Improved Data Structures for LLM Indexing & Response Synthesis
```

The GPT Index project is excellent and worth a read. It utilizes a collection of data structures both created by OpenAI and LangChain:

- List Index - Each node represents a text chunk, otherwise unaltered. In the default setup, all nodes are combined into a single list.
- Vector Store Index - This is equivalent to the simple design that I explained in the previous section. Each text chunk is embedded and stored in a vector store.
- Keyword Index - This supports a quick and efficient lexical search for particular strings.
- Tree Index - This is extremely useful when your data is organized into hierarchies. Consider a clinical document.

GPT Index offers composability of indices, meaning you can build indices on top of other indices. For example, you can combine a Vector Store Index with a Keyword Index to create an Expanded Context Size.

Some of the approaches outlined in this post sound "hacky" because they involve workarounds to the relatively small context size of GPT-4. GPT-4 is anticipated within the next 1-3 months. It is rumored to have a larger context size.

This paper from the folks at Google AI features a number of explorations of engineering tradeoffs. One of the core challenges is how to scale a new state space model architecture scales ~linearly with context size instead of quadratically like seen in traditional models.

In my opinion, advancements in context size will scale alongside demands for more data retrieval; in other words, you need more data to train a model that can handle a larger context size.

Persisting State (e.g. Conversation History)

When LLM's are presented to the user in a conversational form, a major challenge is maintaining that conversation state. This is typically done by persisting the conversation history in a database or file system.

An overview of the relevant strategies is beyond the scope of this post; for an example of a recent code demonstration, see this post.

CHUNKING

Split into a list of paragraphs

```
In [4]: texts = text.split('\n\n')
```

Clean up to remove empty spaces and new lines

```
In [5]: texts = np.array([t.strip(' \n') for t in texts])
```

Checking first 10 paragraphs

```
In [7]: texts[:5]
```

```
Out[7]: array(['This is a fascinating time in the study and application of large language models. New advancements are announced every day!',
               'In this guide, I share my analysis of the current architectural best practices for data-informed language model applications. This particular subdiscipline is experiencing phenomenal research interest even by the standards of large language models - in this guide, I cite 8 research papers and 4 software projects, with a median initial publication date of November 22nd, 2022.',
               'Overview\nIn nearly all practical applications of large language models (LLM's), there are instances in which you want the language model to generate an answer based on specific data, rather than supplying a generic answer based on the model's training set. For example, a company chatbot should be able to reference specific articles on the corporate website, and an analysis tool for lawyers should be able to reference previous filings for the same case. The way in which this external data is introduced is a key design question.',
               'At a high level, there are two primary methods for referencing specific data:',
               'Insert data as context in the model prompt, and direct the response to utilize that information\nFine-tune a model, by providing hundreds or thousands of prompt <> completion pairs\nShortcomings of Knowledge Retrieval for Existing LLM's\nBoth of these methods have significant shortcomings in isolation.'],
              dtype='<U812')
```

Embeddings

```
In [8]: response = co.embed(
        texts=texts.tolist()
    ).embeddings
```

Checking dimension

```
In [11]: embeds = np.array(response)
        embeds.shape
```

```
Out[11]: (36, 4096)
```

Showing the Embeddings(Vector representations)

```
In [12]: embeds
```

```
Out[12]: array([[ 1.8261719 ,  1.3398438 ,  1.9511719 , ...,  0.29418945,
                0.5629883 ,  2.1503906 ],
               [ 1.1611328 , -0.05752563,  0.44995117, ...,  0.6870117 ,
                -1.8222656 ,  0.62402344],
               [-1.2568359 , -0.50097656, -1.1279297 , ...,  1.2841797 ,
                -0.06890869, -0.05160522],
               ...,
               [-1.0722656 , -2.1699219 ,  0.00843048, ...,  0.32739258,
                -1.0058594 , -0.89697266],
               [ 0.11981201, -0.13781738, -0.76953125, ...,  0.26513672,
                -0.6118164 ,  0.42089844],
               [ 0.5991211 ,  0.01847839,  0.5214844 , ...,  0.14221191,
                -0.6118164 ,  0.53808594]])
```

Create the search index

```
In [13]: search_index = AnnoyIndex(embeds.shape[1], 'angular')
# Add all the vectors to the search index
for i in range(len(embeds)):
    search_index.add_item(i, embeds[i])

search_index.build(10) # 10 trees
search_index.save('test.ann')
```

Out[13]: True

Search Function

```
In [18]: pd.set_option('display.max_colwidth', None)

def search(query):

    # Get the query's embedding
    query_embed = co.embed(texts=[query]).embeddings

    # Retrieve the nearest neighbors
    similar_item_ids = search_index.get_nns_by_vector(query_embed[0],1,

                                                    include_distances=True)

    # Format the results
    results = pd.DataFrame(data={'texts': texts[similar_item_ids[0]],
                                'distance': similar_item_ids[1]})

    print(texts[similar_item_ids[0]])

    return results
```

```
In [19]: query = "what are large language models"
search(query)
```

["Language model embeddings are numerical representations of concepts in text and seem to have endless uses. Here's how they work: an embeddings model converts text into a large, scored vector, which can be efficiently compared to other scored vectors to assist with recommendation, classification, and search (+more) tasks. We store the results of this computation into what I'll generically refer to as the search index & entity store - more advanced discussions on that below."]

```
Out[19]:
```

	texts	distance
0	Language model embeddings are numerical representations of concepts in text and seem to have endless uses. Here's how they work: an embeddings model converts text into a large, scored vector, which can be efficiently compared to other scored vectors to assist with recommendation, classification, and search (+more) tasks. We store the results of this computation into what I'll generically refer to as the search index & entity store - more advanced discussions on that below.	0.866105

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js